

SCRUTINIZING COLUMN-ORIENTED DATABASES

OBJECTIVES

This is time to get familiar with the internals of column-oriented databases.

A column-oriented database uses vertical fragmentation to store data. Thus, opposite to row-oriented databases (i.e., traditional relational databases) that store data by rows, they store it by columns. As consequence, in front of random queries (e.g., those of a Data Warehouse) it tends to yield a much better effective read ratio.

In this lab we will scrutinize the differences between row-oriented and column-oriented databases. Although row-oriented database vendors claim they can achieve similar performance with the correct tuning, this turns not to be true.

In this lab you are aimed at implementing the different ways a row-oriented relational database can be tuned for read-only queries. Therefore, you will be asked to tune Oracle 11g in three different ways: using index-only query answering, materialized views and nested tables (to simulates vertical fragmentation) and discuss the query execution carried out by Oracle in each case. You are then asked to compare that execution plan with the internal structures and query processing a column-oriented database does and realize there are some optimizations that a row-oriented will never be able to do.

The list of objectives for this session is as follows:

- Given a read-only workload, decide what of the three main tunings for relational databases give a better result:
 - Index-only query answering techniques (with both B+ and bitmaps),
 - Materialized views,
 - Using UDTs and nested tables to simulate vertical fragmentation.
- Understand the execution plan proposed by Oracle for each case.
- Then, realise about the main drawbacks in these execution plans. Considering how column-oriented databases store and process data, discuss each execution plan with regard to what a column-oriented database would do and identify the differences.

REQUIRED KNOWLEDGE

A script to delete any kind of object, the description of the operations in Oracle's explanation plan and some hints are also provided as additional material.

Your lab mates for this session will be that of the team creation 3 event.

TOOLS

These exercises will be done on Oracle 11g and the Learn-SQL tool.

TRAINING (ACTIVITIES TO DO DURING THE WEEK)

- Solve the training using the Learn-SQL tool about index-only query answering and materialized views (first and second exercise). **To do so, you will need to refresh some basics about Oracle 11g. Check the appendix A for further information.**
- Solve now the same exercise using vertical fragmentation (exercise 3). **To do this part, follow the instructions in appendix B step by step.**
- Which of all three techniques gets a better cost? Think why this is the case. Use the Oracle's execution plan explanation to help yourself. After implementing the needed structures for each solution, execute each query as shown in the following image:

The screenshot shows the Learn-SQL tool interface. At the top, there are tabs for 'abeurp02~2' and 'abeurp02~3'. Below the tabs, there is a toolbar with various icons, including a button labeled 'Rastreo Automático... (F6)' which is circled in red. The main area displays a SQL script with several queries. The first query is highlighted in blue. Below the script, there is a table showing the execution plan for the selected query. The table has four columns: OPERATION, OBJECT_NAME, COST, and LAST_CR_BUFFER_GETS. The first row shows 'SELECT STATEMENT' with a cost of 36. The second row shows 'MAT_VIEW REWRITE ACCESS (FULL)' with a cost of 4.

OPERATION	OBJECT_NAME	COST	LAST_CR_BUFFER_GETS
SELECT STATEMENT		36	
MAT_VIEW REWRITE ACCESS (FULL)	V3	4	36

Zoom the figure in, and you will see the “rastreo automático F6” button (which is circled). When you select a query (for example, one of those in the statement) and press this button, Oracle executes the query (be careful, it is not an estimation but the real execution) and shows the execution plan.

In this example, it is using a materialized view (V3) to rewrite the query shadowed in blue in terms of this view and therefore, although it estimates 36

blocks of cost, the real cost is 4 (LAST_CR_BUFFER_GETS is the number of cache reads –i.e., blocks- needed to answer the query during the execution).

- Are you able to understand each of the plans? Do they make sense to you? Feel free to make notes and ask the lecturer during the lab if you still have doubts. You can check a detailed description of each operation in Oracle's execution plan in the additional material of this session.

Remember you are responsible for understanding the main objectives of this session beyond the training exercises. Thus, if you need further training go for it. Finally, bear in mind these exercises have *many* different solutions. If you think Learn-SQL is misbehaving in front of a potentially correct solution, please, contact the lecturer.

DELIVERABLES

To deliver these exercises you must set your Oracle 11g account with your solution and provide your login and password in the Learn-SQL tool. In the lab you will be asked to solve exercises similar to these and discuss if the execution plan of your queries could be outperformed by a column-oriented database.

OPTIONAL ACTIVITIES

Repeat these exercises in a column-oriented database:

- Go to <http://www.monetdb.org/Home> and download MonetDB, an open-source column-oriented database.
- Implement the schema proposed in the training and launch the queries.
- Ask MonetDB to let you know the execution plan. This is a real column-oriented database! You will see the differences with Oracle.

- Read the lab instructions.

- Get familiar with your Oracle account. If you want to connect from home you will need to use a VPN client to connect remotely to an Inlab machine. Remember: if you are running Windows you will need to execute the VPN client in administrator mode or it will not get the needed grants to connect.

- Before going to the lab, clean your Oracle account properly (you will gain time, remember that during the session, with all your peers there, it will be slower than at home).

- Be sure to completely clean your Oracle account after each try. You can use the SQL script provided (see lab material) to do so but bear in mind that it might happen some objects remain after executing it (e.g., objects with dependences, additional objects, etc.) To be sure your account is completely empty execute the following query: `SELECT * FROM USER_TS_QUOTAS;`

- Follow the order set in the attachment files. There is a clear flow you can follow to avoid problems (e.g., not updating the statistics or forgetting doing some steps). If you interrupt an execution, it is better to start from scratch to avoid weird situations.

- To know if a bitmap has more than 100 repetitions for each value you can do as follow: number of tuples in the indexed table / distinct keys in the bitmap.

- For exercise 3, refresh the nested tables notation to adapt the inserts scripts and needed views for correction mentioned in the instructions. Remember the Oracle manual is available here.

- After each try, check the access plan and try to draw your own conclusions. In the additional material of this lab you will find a detailed description of each operation in Oracle's explanation plan.

- If you want to execute some of the statements in the attachments at once, just add / BELOW (not after, not at the side, below in a standalone line) each execution block (i.e., statements ended in ; or after the END keyword of a PL/SQL block).

APPENDIX A

A detailed step by step guide for exercises 1 and 2 in the training

Each Learn-SQL exercise has an attachment. There, you will find the following parts:

- Database schema (of the exercise): It contains a set of CREATE tables. Remember tables must be created with PCTFREE 0 and ENABLE ROW MOVEMENT as explained later in this appendix.
- Inserts: Data you must insert in the tables. Just execute this script. Note that, at the end, you must execute "ALTER TABLE *table_name* SHRINK SPACE". This guarantees the table is compacted.
- Next, you are supposed to insert your solution (i.e., indexes and / or materialized views).
- The next code updates the catalog statistics (needed to keep the database catalog up to date).
- Finally, execute the last part, creating the measure table, to know the average cost for your solution. Note that each query is properly weighted according to the weights given in the statement. Some insights about this code:
- The query cost taken into account is that provided by Oracle in the execution (measured in terms of gets from the cache) for each query (weighted by the query).
- The used space value you must honor, if there is a space constraint, is the one corresponding to USER_TS_QUOTAS. After creating your structures and inserting the needed data, trigger this query to check it:

```
SELECT SUM(BLOCKS) FROM USER_TS_QUOTAS
```

CONSTRAINTS

These constraints are needed to let Learn-SQL properly assess your solution. These are rules about how to create the structures, which have been created in the following way and no other:

- Table without clustered index, nor clustered structure, nor hash:

```
CREATE TABLE name (attributes) PCTFREE 0 ENABLE ROW MOVEMENT;
```

Insertions

```
ALTER TABLE name SHRINK SPACE; // Compress the table
```

- B+ tree (after the insertions):

```
CREATE TABLE name (attributes) PCTFREE 0 ENABLE ROW MOVEMENT;
```

Insertions

```
ALTER TABLE name SHRINK SPACE; // Compress the table
CREATE INDEX name ON table (attributes) PCTFREE 33;
```

- We create the index after the insertions to assure that it holds the desired load factor (normally, 2/3).

- Bitmap¹:

```
CREATE TABLE name (attributes) PCTFREE 0 ENABLE ROW MOVEMENT;
```

Insertions

```
ALTER TABLE name SHRINK SPACE; // Compress the table
```

```
ALTER TABLE table MINIMIZE RECORDS_PER_BLOCK; // Compress
the bitmap
```

```
CREATE BITMAP INDEX name ON table(attributes) PCTFREE 0;
```

- **Bitmaps should only be considered for attributes with values repeated – over 100 times. To know how many times is repeated each value of a bitmap compute it as follows: #rows in the table / # distinct keys in the index (check the catalog, USER_TABLES and USER_INDEXES to find these values).**

- Materialized views:

```
CREATE MATERIALIZED VIEW view_name ORGANIZATION HEAP PCTFREE 0
BUILD IMMEDIATE REFRESH COMPLETE ON DEMAND ENABLE QUERY REWRITE AS
( query );
```

- **Do not consider FAST ON COMMIT views.**

General rules if you want to go for any other solution (although the maximum mark can be obtained just considering the above elements):

¹ Due to their compression techniques, bitmaps generate little unforeseen variations in query costs (usually one or two blocks). This potential deviation will never result in more than one point in the exercise, and will be manually (off-line) corrected by the proctor.

- “ROW MOVEMENT” must be enabled for any table not in a cluster. It will allow you to “SHRINK SPACE” just after the insertions.
- “MINIMIZE RECORDS_PER_BLOCK” must be executed after the insertions of the table and before the creation of any bitmap index on it (you only need to execute it once, even if you create many bitmap indexes over the same table).

ADDITIONAL INFORMATION

- You might be interested in the following catalog tables. They will let you know what is happening and give the actual values for space and cost (inside parenthesis you have some attributes that may be interesting):
 - USER_TABLES (TABLE_NAME, CLUSTER_NAME, IOT_TYPE, IOT_NAME, PCT_FREE, BLOCKS, NUM_ROWS, AVG_ROW_LEN, LAST_ANALYZED)
 - USER_TAB_COLS (TABLE_NAME, COLUMN_NAME, DATA_TYPE, DATA_LENGTH, AVG_COL_LEN, NULLABLE, LAST_ANALYZED)
 - USER_INDEXES (INDEX_NAME, TABLE_NAME, INDEX_TYPE, UNIQUENESS, PCT_FREE, BLEVEL, LEAF_BLOCKS, DISTINCT_KEYS, LAST_ANALYZED, JOIN_INDEX)
 - USER_SEGMENTS (SEGMENT_NAME, SEGMENT_TYPE, BYTES, BLOCKS)
 - USER_TS_QUOTAS (TABLESPACE_NAME, BYTES, BLOCKS)
 - RECYCLEBIN (ORIGINAL_NAME)

APPENDIX B

A detailed step by step guide for exercise 3 in the training

In this exercise you are asked to use vertical fragmentation to solve the same statement as in exercise 1 and 2. However, using object-relational features changes a bit how to proceed. To avoid getting lost, follow these steps to modify the attached file and build your solution:

- Decide how to vertically fragment the database. You may want to use the affinity matrix method we saw in the lectures. For each vertical fragment you must create a nested table. Here, you need to assess whether a pure vertical partitioning (i.e., each column is stored separately in a nested table) or a hybrid approach (i.e., group some columns in each nested table).
- If you used the affinity matrix method, be careful; Oracle's best solution could not be the one theoretically identified. Do you see why? Oracle nested tables are, in the end, row-oriented, thus, there is extra costs hidden that spoil the result.
- Replace the create table with the corresponding CREATEs for the collection types and the table containing the nested tables. As before, any table must be created with PCTFREE 0 and ENABLE ROW MOVEMENT (this does not apply to nested tables).
- Modify the INSERT script to adapt it to your new object-relational schema (basically, you need to insert data into nested tables now and thus you may want to refresh the previous session on object-relational).
- Finally, to adapt the last part, you need to create three views view1, view2 and view3. Each view must adapt each of the queries in the exercise statement to your new object-relational schema. For example, consider Q1:

```
SELECT cand AS a, AVG(val) AS b FROM poll_answers  
GROUP BY cand);
```

You must create a view view1 adapting this query to your nested tables and retrieving the same data. For example, assume cand and val are in two different nested tables p1 and p2 then (how to query nested tables is, again, material from the previous lab):

```
CREATE VIEW view1 AS  
SELECT p1.cand AS a, AVG(p2.val) AS b  
FROM poll_answer pa, TABLE(pa.part1) p1, TABLE(pa.part2) p2  
GROUP BY p1.cand
```

IMPORTANT: Each attribute in the view select must have an alias, and these aliases must coincide with the alphabet sorted lexicographically. In other words, the first attribute must have a as alias (see the example), the

second one b, the third c, and so on. Failing to do so, your exercise will not be correctly assessed by Learn-SQL.

All the rest remain the same and you can still ask Oracle for the execution plan of a query. Remain to update the statistics before sending your answer.